

Enterprise AI Awareness Training Series

Responsible AI & Ethics

Module 2 of 5 | Duration: 45 Minutes

Audience: All Employees — No Technical Background Required

Table of Contents

Note: After opening this document in Microsoft Word, right-click the Table of Contents and select “Update Field” to refresh page numbers.

1. Course Overview

Purpose of This Training

AI tools are now woven into hiring, customer communication, marketing, planning, and daily writing. With that reach comes responsibility: AI can amplify bias, produce convincing falsehoods, blur accountability, and create legal exposure — quietly and at scale. This module equips every employee to use AI in a way that is fair, transparent, and defensible.

Responsible AI is not a policy document that lives with the legal team. It is a set of everyday judgements made by the people who actually use the tools: which task to give the AI, what to check before relying on output, when to disclose AI involvement, and when to escalate a concern.

Learning Objectives

By the end of this 45-minute module, you will be able to:

- Explain the core principles of responsible AI: fairness, transparency, explainability, accountability, and human oversight.
- Recognize how bias enters AI systems and spot warning signs in everyday outputs.
- Apply a simple ethical decision-making framework before using AI for a sensitive task.
- Identify misuse risks — including deepfakes, impersonation, and deceptive content — and respond appropriately.
- Understand copyright and intellectual-property considerations when using AI-generated content.
- Know when and how to disclose AI use, and when to escalate concerns.

Why This Topic Matters in Today's Workplace

- **Scale:** A biased human decision affects one case; a biased AI-assisted process can affect thousands of cases before anyone notices.
- **Trust:** Customers, candidates, and colleagues lose trust quickly when AI is used deceptively or carelessly — and trust is far harder to rebuild than to keep.
- **Accountability:** “The AI did it” is not a defense. Legally and professionally, responsibility for AI-assisted work stays with the people and organization that use it.
- **Regulation:** AI-specific laws and regulations are emerging worldwide, with particular focus on high-impact uses such as hiring, credit, and biometric identification. Organizations are expected to demonstrate responsible practices.

The one-sentence version

If you use AI to help produce a decision, document, or message, you own the outcome — so use approved tools, check for fairness and accuracy, be honest about AI's role, and keep a human meaningfully in charge.

2. Introduction

What “Responsible AI” Means in Simple Terms

Responsible AI means using AI in ways that are fair to people, honest about what the AI is and does, and accountable — meaning a human or organization clearly answers for the results. It applies at two levels:

- **Organizational level:** choosing, configuring, and governing AI systems carefully — vendor checks, testing for bias, policies, monitoring.
- **Individual level:** the daily choices each employee makes about what to ask AI to do, what to check, and what to disclose. This module focuses here.

Business Relevance

Ethical AI use is not a brake on productivity — it is what makes AI adoption sustainable. Organizations that skip it tend to discover the cost later: discrimination complaints arising from AI-screened hiring, public embarrassment from fabricated content, copyright disputes over AI-generated marketing, or customer backlash after undisclosed chatbot interactions. Each of these is far more expensive than the review steps that would have prevented it.

Common Misconceptions

Misconception	Reality
“AI is neutral because it's just math.”	AI learns from human-generated data, which carries human biases. Math faithfully reproduces the patterns it is given — including unfair ones.
“Ethics is the legal team's job.”	Policies set boundaries, but most ethical decisions happen at the keyboard: what you ask, what you check, what you send.
“If the tool allows it, it must be acceptable.”	Tools rarely enforce ethics. A tool that can generate a fake testimonial will do so; whether you should is your judgement.
“Responsible AI means avoiding AI.”	It means using AI deliberately: right task, right oversight, honest disclosure. Refusing all AI use is not more ethical — it just forgoes the benefits.
“Bias only matters in hiring.”	Bias can surface in marketing imagery, customer prioritization, performance summaries, translation tone, and more.

“If I didn't write it, I'm not responsible for it.”

You are responsible for everything you submit, send, or publish — regardless of what drafted it.

3. Key Concepts

3.1 Responsible AI Principles

Most organizational and governmental frameworks converge on a similar set of principles. Here they are with their everyday meaning:

Principle	Plain meaning	What it looks like in daily work
Fairness	AI should not systematically disadvantage people or groups	Reviewing AI-assisted screening or ranking for skewed patterns
Transparency	Be honest that AI is involved	Labelling chatbots; disclosing AI assistance where policy requires
Explainability	Decisions should be understandable, not a black box	Being able to say why a candidate or claim was prioritized
Accountability	A named human/organization answers for outcomes	You sign off on AI-assisted work as your own
Human oversight	Humans stay meaningfully in control of consequential decisions	AI recommends; a person decides, with real power to overrule
Privacy & safety	Protect people's data and wellbeing	Covered in depth in Module 3

3.2 Fairness and Bias

Bias in AI means systematic unfairness in outputs — favoring or disadvantaging certain groups. It usually enters through training data: if historical data reflects past discrimination or under-representation, the model learns those patterns as “normal.”

A widely cited real-world example: a major technology company scrapped an experimental CV-screening tool after discovering it penalized résumés associated with women — because it had learned from a decade of male-dominated hiring data. The system was not programmed to discriminate; it learned to.

Bias can be subtle in everyday outputs, for example:

- Image generation that depicts “a CEO” or “a nurse” with a consistent gender or ethnicity.
- Drafted job adverts whose wording skews masculine or feminine in tone.
- Summaries that emphasize some speakers' contributions over others.
- Customer-prioritization suggestions that correlate with postcode or name patterns.

Practical habit

When AI output involves people — candidates, customers, employees — pause and ask: “Would this output look different if the people involved were a different gender, ethnicity, age, or background?” If you're unsure, get a second pair of eyes.

3.3 Transparency

Transparency means not hiding AI's involvement from people who have a right to know. Three common workplace forms:

- **Disclosure to recipients:** customers should know when they are talking to a chatbot rather than a human.
- **Disclosure within teams:** telling your manager or collaborators when significant work was AI-assisted, where policy or professional norms require it.
- **Disclosure in published content:** following your organization's rules on labelling AI-generated images, text, or media.

Transparency does not mean footnoting every spell-check. Use judgement and policy: the more the recipient's trust or decision depends on the content's origin, the stronger the duty to disclose.

3.4 Explainability

Explainability is the ability to give an understandable reason for an AI-influenced decision. It matters most when decisions affect people: hiring, promotion, credit, claims, discipline. “The system scored them lower” is not an explanation; it is an admission that no one is in control.

Workplace example: If an AI tool helps shortlist candidates, the recruiter should be able to state the job-relevant reasons a candidate did or did not progress — and those reasons should survive scrutiny without referencing “the algorithm”.

3.5 Accountability

Accountability answers one question: who is responsible when AI-assisted work is wrong, unfair, or harmful? The answer is never “the AI.” Software cannot be accountable; it cannot explain itself to a regulator, apologize to a customer, or learn a professional lesson.

Accountability stays with:

- The employee who used the AI and submitted the work.
- The manager or function that approved the process.
- The organization that deployed the tool.

Memorize this

You can delegate the typing to AI. You cannot delegate the responsibility.

3.6 Human Oversight

Human oversight means a person remains meaningfully in the loop for consequential outcomes — not as a rubber stamp, but with the time, information, and authority to disagree with the AI. Two failure modes to avoid:

- **Automation bias:** humans tend to over-trust machine suggestions, approving them faster than they would a colleague's. A reviewer who clicks “accept” 200 times an hour is not oversight.
- **Rubber-stamp design:** processes that technically include a human but give them no realistic way to check. Oversight must be designed to be possible.

Rule of thumb for matching oversight to stakes:

Stakes of the task	Appropriate oversight
Low (internal brainstorm, rough draft)	Quick read-through by the user
Medium (client email, report, published copy)	Full review and fact-check by the user; normal approval chain
High (decisions about people, legal/financial commitments, public statements)	AI may inform but not decide; documented human decision; second reviewer where required

3.7 Ethical Decision Making: A Quick Framework

Before using AI for a task that involves people, money, or public communication, run the PACT check:

Step	Question to ask yourself
P — Permission	Is this use allowed by our AI policy, and is the tool approved? Is the data appropriate to share with it? (See Module 3.)
A — Accuracy & fairness	Have I verified the output, and could it be biased against anyone affected?
C — Candor	Would I be comfortable if everyone affected knew exactly how AI was used here? If not, why not?
T — Takeover	Is a human still meaningfully making the consequential decision, with authority to overrule?

If any answer is “no” or “not sure,” pause and either adjust the approach or ask — your manager, HR, legal, or whoever owns the AI policy.

3.8 Misuse of AI

The same capabilities that make AI useful make it abusable. Forms of misuse employees should never engage in — and should report if they encounter:

- **Deception:** fake reviews or testimonials, fabricated case studies or statistics, AI-written messages impersonating a real person without authorization.
- **Manipulation:** generating deliberately misleading content, astroturfing, or exploiting psychological pressure points at scale.
- **Harassment or defamation:** using AI to mock, demean, or fabricate claims about colleagues, competitors, or anyone else.
- **Academic/professional dishonesty:** presenting AI-generated analysis as independently performed expert work where authenticity is expected (e.g., certifications, audits, signed opinions).
- **Circumvention:** trying to trick AI tools into bypassing their safety rules, or using unapproved tools to dodge company controls.

3.9 Deepfakes

Deepfakes are AI-generated or AI-altered audio, video, or images that convincingly depict real people saying or doing things they never did. Two workplace angles:

- **Never create them:** synthesizing a colleague's, executive's, or customer's voice, face, or likeness without explicit authorization is off-limits — even “as a joke.” Legitimate uses (e.g., an approved avatar of a consenting executive for training videos) go through formal approval.
- **Stay alert to receiving them:** criminals have used cloned executive voices and even deepfaked video calls to order urgent money transfers. In one widely reported case, an employee transferred tens of millions of dollars after a video conference in which every other participant was a deepfake. Treat unusual or urgent requests — especially involving payments, credentials, or confidential data — as unverified until confirmed through a separate, known channel (call back on a saved number; confirm in person).

Red flags for synthetic media

Urgency plus secrecy (“do this now, tell no one”), requests that bypass normal process, slightly “off” voice or video quality, refusal to switch channels, or a request the real person would be unlikely to make. When in doubt, verify out-of-band and report it.

3.10 Copyright Considerations

Three questions cover most everyday situations. (Note: this is general awareness, not legal advice — route real disputes to your legal team.)

- **1. What goes in:** Don't feed AI tools content you don't have the right to use — e.g., pasting a full paid industry report or a competitor's copyrighted material to repurpose it.
- **2. What comes out:** AI output can occasionally resemble existing copyrighted works, brands, or distinctive styles. For published or commercial content — slogans, images, music, long text — run normal originality and brand checks before use.

- **3. Who owns it:** Ownership of AI-generated content varies by jurisdiction and tool terms; in several jurisdictions, purely AI-generated works may have weak or no copyright protection. For brand-critical assets, involve legal before relying on exclusivity.
- **Also:** never use AI to reproduce song lyrics, book passages, paid content, or other protected material for distribution — “the AI wrote it out” does not make copying lawful.

3.11 Workplace AI Ethics in One Table

Situation	Ethical anchor
AI helping with decisions about people (hiring, reviews, discipline)	Human decides; reasons must be explainable and job-relevant; check for bias
AI drafting external communications	Verify facts; recipient deception is off-limits; disclose where required
AI-generated media (images, audio, video)	No unauthorized real-person likenesses; label per policy; originality check
AI summarizing colleagues' words (meetings, feedback)	Check fairness of representation; sensitive feedback needs source review
Receiving suspicious AI-era content (voice notes, urgent requests)	Verify out-of-band; report suspected deepfakes or fraud immediately
Feeling pressure to over-rely on AI to hit deadlines	Raise it; speed never justifies skipping review on consequential work

4. Scenario-Based Learning

Nine scenarios show responsible-AI judgement in action across functions. Read the situation, decide what you would do, then compare.

Scenario 1 (HR): The AI-Ranked Shortlist

Situation: Nadia in HR pilots an AI tool that ranks 300 applications for a finance role. Reviewing the top 30, she notices almost all come from the same two universities and a narrow age band, although strong candidates from other backgrounds appear further down the list.

Challenge: The ranking may reflect patterns in historical hiring data rather than job-relevant merit — a fairness risk with legal implications.

Recommended approach: Nadia doesn't simply accept the top 30. She flags the pattern to the hiring manager and the tool's owner, samples candidates from across the full list to compare against job criteria, and documents the concern. The team treats AI ranking as one input, with humans applying the actual criteria — and asks the vendor how the tool was tested for bias.

Explanation: Skewed outputs that correlate with background characteristics are the classic symptom of learned bias. Fairness requires noticing the pattern, not assuming the machine is objective. Explainability requires being able to defend each shortlisting decision on job-relevant grounds.

Key takeaway

In decisions about people, AI suggests — humans decide, on explainable, job-relevant grounds.

Scenario 2 (Sales): The Tempting Fake Testimonial

Situation: Jonas is building a proposal and the deadline is tonight. A colleague suggests: “Just have the AI write a couple of customer quotes — make them sound like our real clients, nobody checks.”

Challenge: Fabricated testimonials are deceptive, likely violate marketing law and company policy, and could destroy a client relationship if discovered.

Recommended approach: Jonas refuses. Legitimate alternatives exist: ask AI to draft an outreach note requesting real testimonials, summarize genuine feedback already on file (with permission), or present anonymized real results. He also pushes back on the colleague — and if the practice seems common, raises it with his manager.

Explanation: This is misuse by deception. The ethical line isn't whether anyone checks; it's whether the content misrepresents reality to influence a decision. AI lowers the effort of fabrication, which is precisely why personal integrity and team norms matter more, not less.

Key takeaway

AI makes fabrication easy; that's a reason for firmer lines, not looser ones.

Scenario 3 (Marketing): The Suspiciously Familiar Campaign Image

Situation: Mei generates campaign images and one option is striking — but it strongly resembles the visual style and composition of a famous photographer's well-known series. It would look great on the billboard.

Challenge: AI-generated content can echo existing copyrighted works or distinctive styles closely enough to create legal and reputational risk in commercial use.

Recommended approach: Mei treats resemblance as a stop sign for commercial use. She runs the image past brand/legal per the normal creative-approval process, uses reverse-image search to check similarity to existing works, and generates alternatives with more original direction. For the billboard, originality review is mandatory, not optional.

Explanation: The “what comes out” copyright question in action. The convenience of generation doesn't change commercial-use standards: published assets must clear the same originality and rights checks as agency-produced work.

Key takeaway

AI-generated does not mean rights-cleared. Commercial use requires the normal checks.

Scenario 4 (Operations): The Voice on the Phone

Situation: Late Friday, an operations coordinator receives a voice message that sounds exactly like the COO: “We're closing a confidential acquisition. Wire the deposit to this account within the hour and don't discuss it with anyone.” The number is unfamiliar; the voice is convincing.

Challenge: Voice cloning makes “I recognized the voice” obsolete as verification. Urgency plus secrecy is the signature of fraud.

Recommended approach: She does not act on the message. She verifies through a separate, known channel — calling the COO's saved number or contacting their office directly — and reports the message to security/IT immediately, regardless of outcome. No legitimate executive request requires bypassing payment controls and secrecy from colleagues.

Explanation: Deepfake-enabled fraud succeeds by hijacking authority and urgency. The defense is procedural, not perceptual: out-of-band verification and fixed payment controls that no phone call can override.

Key takeaway

Voices and videos can be faked. Processes and call-backs cannot — verify out-of-band, then report.

Scenario 5 (Project Management): The Lopsided Meeting Summary

Situation: Omar uses AI to summarize a tense project retrospective. Reading the draft, he notices it attributes most constructive ideas to two senior voices and renders a junior engineer's criticism — which actually drove the meeting's key decision — as a vague “some concerns were raised.”

Challenge: Summaries shape reputations and records. Uneven representation, even unintentional, can quietly disadvantage quieter or junior contributors.

Recommended approach: Omar corrects the summary against his own notes before circulating, explicitly crediting the engineer's contribution. Going forward, he prompts the tool to attribute points to speakers where the transcript supports it, and gives participants a chance to correct the record.

Explanation: Fairness isn't only about hiring algorithms; it shows up in everyday artifacts like minutes and feedback summaries. The human circulating the document is accountable for what it says about people.

Key takeaway

When AI summarizes people's words, check that credit and criticism land fairly before you hit send.

Scenario 6 (Customer Support): The Chatbot That Pretends to Be Human

Situation: A support team lead proposes renaming the customer-facing chatbot “Sarah,” giving it a human photo, and instructing it to deny being a bot if asked, reasoning that “customers trust humans more.”

Challenge: Deliberately deceiving customers about whether they're talking to AI violates transparency norms, several jurisdictions' rules, and customer trust.

Recommended approach: The team pushes back: the bot can be friendly and named, but it must be identified as a virtual assistant, must never claim to be human, and must offer a clear path to a real person. The trust problem is solved by making the bot genuinely useful and escalation easy — not by lying.

Explanation: Transparency to recipients is the principle at stake. Short-term metrics gained through deception convert into long-term damage the moment customers discover it — and they do.

Key takeaway

Bots may be friendly; they may not be undercover.

Scenario 7 (Leadership): “The Algorithm Decided”

Situation: A department head, Yusuf, is asked by an employee why their internal transfer request was declined. The process used an AI-based fit score. The HR partner's file contains only the score: 41/100, below the 60 threshold. No one can say what drove the number.

Challenge: A consequential decision about a person has been made with no explainable reasons — indefensible to the employee, and risky if challenged formally.

Recommended approach: Yusuf declines to defend the unexplained score. He has the request re-reviewed by humans against stated criteria and gives the employee genuine reasons. He then escalates the process gap: either the tool's owners provide explainable, validated criteria and human decision rights, or the tool shouldn't gate people decisions.

Explanation: Explainability and accountability are inseparable: if no one can explain a decision, no one is genuinely accountable for it. “The algorithm decided” signals that human oversight existed on paper only.

Key takeaway

If you can't explain a people decision without citing a score, the decision isn't ready to stand.

Scenario 8 (Finance): The Polished Paragraph with Invented Authority

Situation: Preparing a budget memo, Carmen asks AI to strengthen her argument. The output adds: “Industry benchmarks indicate companies of our size allocate 12–15% of revenue to this function.” It sounds authoritative. She has no source for it.

Challenge: Inserting unverified, possibly fabricated 'benchmarks' into a decision document misleads leadership — even if the writer didn't intend deception.

Recommended approach: Carmen deletes or verifies the claim before sending. If she can find a real, citable benchmark, she cites it; if not, she argues from internal data she can stand behind. She treats every authoritative-sounding AI addition — statistics, studies, expert consensus — as unsourced until proven otherwise.

Explanation: Honesty in internal documents is an ethics issue, not just an accuracy issue: decision-makers are entitled to know which claims rest on evidence. Hallucinated authority is a subtle form of deception that the sender owns once they hit send.

Key takeaway

If you can't source it, don't send it. AI-added 'facts' are claims you are now making.

Scenario 9 (General Office): The Parody Video of a Colleague

Situation: For a team offsite, someone proposes using an AI video tool to make a humorous clip of a well-liked manager 'singing' a pop song, using photos and a voice clone from meeting recordings. "He'll love it — it's affectionate."

Challenge: Creating synthetic media of a real colleague without consent — even affectionately — crosses a line on likeness, consent, and precedent, and the clip can escape its context.

Recommended approach: The team doesn't create it without asking. The minimum bar is the manager's informed consent in advance; even then, company policy on synthetic media of employees applies, and anything using cloned voice from internal recordings deserves a hard pause. Safer alternatives: a clearly fictional animated character, or getting the manager to film a real cameo.

Explanation: Consent is the dividing line for real-person likenesses, and 'positive intent' doesn't substitute for it. Normalizing casual deepfakes of colleagues also erodes the team's ability to treat synthetic media of executives as a red flag — the security habit from Scenario 4.

Key takeaway

No synthetic likenesses of real people without informed consent and policy approval — humor included.

5. Best Practices

Dos and Don'ts

Do	Don't
Keep humans meaningfully in charge of decisions about people, money, and public commitments.	Don't let an AI score or ranking substitute for explainable human judgement.
Check outputs involving people for skew or unfair representation.	Don't assume "the machine is objective" — it learned from human data.
Be transparent: identify bots as bots; disclose AI assistance where policy or trust requires.	Don't pass off AI work as independently verified expertise, or let a bot claim to be human.
Verify authoritative-sounding claims, statistics, and citations before sending.	Don't let hallucinated "benchmarks" or "studies" into decision documents.
Get informed consent and policy approval before any use of a real person's likeness or voice.	Don't create synthetic media of colleagues, executives, or customers — even as a joke.
Verify unusual or urgent requests out-of-band, and report suspected deepfakes or fraud.	Don't act on urgency-plus-secrecy requests, however convincing the voice or video.
Run originality and brand checks on AI-generated content before commercial use.	Don't publish AI output that may echo existing works, and don't feed in content you lack rights to.
Use the PACT check (Permission, Accuracy & fairness, Candor, Takeover) for sensitive tasks.	Don't proceed on "the tool allowed it" as your ethical standard.

Common Mistakes to Avoid

- **Objectivity assumption:** treating AI rankings of people as neutral because they're numerical.
- **Diffused responsibility:** everyone assumes someone else checked; no one did. Name the reviewer.
- **Rubber-stamp oversight:** a human approval step too fast or uninformed to ever say no.
- **Intent excuse:** "I didn't mean to mislead" — recipients experience the deception either way; verify before sending.

- **Silent escalation gap:** noticing a biased output or shady practice and saying nothing. Reporting concerns is part of responsible use.

Recommended Workplace Behavior

- Apply review effort in proportion to impact on people: the more a task affects someone's opportunities, money, or reputation, the more scrutiny it needs.
- When you spot bias, fabrication, or misuse, raise it — with your manager, the tool owner, or the channel your organization provides. Quietly working around problems leaves them in place for others.
- Be the colleague who asks “how would we explain this decision?” and “would we be comfortable if this were disclosed?” — those two questions catch most ethical problems early.

6. Summary

Key Learnings

- Responsible AI rests on fairness, transparency, explainability, accountability, and human oversight — practiced daily by users, not just written in policy.
- Bias enters through training data and surfaces subtly; outputs about people deserve a fairness check.
- Accountability cannot be delegated to software: you own what you submit, send, or publish.
- Deception is the bright line in misuse — fake testimonials, undercover bots, fabricated authority, and unauthorized deepfakes are all variants of the same violation.
- Deepfakes cut both ways: never create unauthorized likenesses, and never trust urgent requests on voice/video recognition alone.
- AI-generated content needs the same originality, rights, and approval checks as any other content before commercial use.

Important Reminders

- “The algorithm decided” is never an acceptable explanation for a decision about a person.
- Urgency + secrecy + payment or credentials = verify out-of-band, then report.
- When in doubt, run the PACT check — and if still in doubt, ask before acting.

Action Points

- Identify one AI-assisted task in your role that affects other people, and define what your fairness/accuracy check looks like for it.
- Find out how your organization wants suspected deepfakes, fraud attempts, or AI misuse reported — know the channel before you need it.
- Review your team's customer-facing AI touchpoints (if any) for honest labelling.
- Complete the Knowledge Assessment Quiz and Final Assessment that follow.

7. Knowledge Assessment Quiz

20 questions. Question types: Multiple Choice (one answer), Multiple Select (one or more answers), Situational Judgement, and True/False. The questions test application of concepts, not memorization. Answers and explanations follow the quiz.

Question 1 (MCQ)

An AI résumé-screening tool consistently ranks graduates of two universities at the top, mirroring the company's historical hiring. A recruiter says: "It's data-driven, so it's objective." The most accurate response is:

- A. Correct — numerical scores cannot be biased.
- B. Incorrect — the tool learned patterns from historical hiring data, so it can reproduce past bias with a veneer of objectivity; the pattern warrants investigation.
- C. Incorrect — AI tools are forbidden in recruitment.
- D. Correct — universities are a job-relevant criterion, so any ranking based on them is fair.

Question 2 (Situational Judgement)

You receive a voice message that sounds exactly like your CFO: "Urgent and confidential — process this payment in the next 30 minutes, don't loop anyone in." Your best first action is:

- A. Process the payment; you clearly recognized the voice.
- B. Reply to the same number asking for written confirmation.
- C. Verify through a separate, known channel (e.g., call the CFO's saved number or their office), do not act until confirmed, and report the message to security regardless of outcome.
- D. Forward the request to a colleague so the decision is shared.

Question 3 (True/False)

True or False: If an AI tool produced a discriminatory shortlist, the organization and the people who relied on it can still be held responsible, even though no human intended to discriminate.

- A. True
- B. False

Question 4 (MCQ)

A teammate suggests instructing your customer chatbot to claim it's a human agent named "Sarah" because "customers open up more to people." Which principle does this most directly violate?

- A. Explainability

- B. Transparency — recipients are entitled to know they're interacting with AI; deception about this damages trust and may breach regulations.
- C. Tokens
- D. Fine-tuning

Question 5 (Multiple Select)

Which of the following constitute misuse of AI in the workplace? (Select all that apply.)

- A. Generating fake customer testimonials for a proposal
- B. Using AI to brainstorm genuine interview questions for candidates
- C. Creating a voice clone of a colleague for an office joke without their consent
- D. Asking AI to draft a polite payment-reminder email you then review
- E. Inserting an AI-invented “industry benchmark” statistic into a budget memo without verification

Question 6 (MCQ)

An employee is told their transfer was declined because “the AI fit score was 41, below the 60 cutoff,” and no one can say what drove the score. What is the core problem?

- A. The cutoff should have been 50.
- B. The decision lacks explainability — no one can give job-relevant reasons — which also means no one is genuinely accountable for it.
- C. The employee should have asked the AI directly.
- D. Scores should be hidden from employees to avoid such questions.

Question 7 (True/False)

True or False: Content generated by AI is automatically free of copyright concerns because no human author copied anything.

- A. True
- B. False

Question 8 (Situational Judgement)

Reviewing an AI-drafted summary of a team retrospective, you notice a junior colleague's pivotal criticism has been reduced to “some concerns were raised,” while senior voices are quoted by name. Before circulating, you should:

- A. Send it as-is; summaries are inherently lossy.
- B. Delete the criticism entirely to keep the summary positive.

- C. Correct the summary against your notes so contributions are attributed fairly, and credit the colleague's point explicitly.
- D. Ask the AI to make the summary shorter.

Question 9 (MCQ)

What is “automation bias”?

- A. A model's preference for automating manual tasks.
- B. The human tendency to over-trust machine-generated suggestions and approve them with less scrutiny than a colleague's work would receive.
- C. Bias that only occurs in robotic process automation.
- D. An AI's tendency to favor automated industries in its answers.

Question 10 (MCQ)

In the PACT ethical check, the “Candor” question is best expressed as:

- A. “Is the tool's interface candid and user-friendly?”
- B. “Would I be comfortable if everyone affected knew exactly how AI was used here?”
- C. “Has the model been fine-tuned recently?”
- D. “Is the output written in a candid tone?”

Question 11 (Multiple Select)

Which are legitimate, responsible uses of AI involving real people? (Select all that apply.)

- A. An approved training video using an executive's avatar created with their informed consent through the formal process
- B. A humorous deepfake of a manager singing, made secretly as a surprise
- C. AI-suggested interview scheduling emails reviewed before sending
- D. A chatbot that identifies itself as a virtual assistant and offers escalation to a human
- E. Cloning a former employee's voice for the IVR system because callers liked it

Question 12 (Situational Judgement)

AI added this line to your draft memo: “Studies show 73% of firms in our sector have adopted this approach.” You have no source. The deadline is in an hour. You should:

- A. Keep it — it strengthens the argument and is probably roughly right.
- B. Keep it but change “Studies show” to “It is widely known that.”
- C. Verify it against a real source and cite it, or remove it and argue from evidence you can stand behind.

D. Keep it and add “(source: AI)” as the citation.

Question 13 (MCQ)

Which situation is the strongest signal of a possible deepfake-enabled fraud attempt?

- A. A routine calendar invitation from a known colleague.
- B. An unusual video-call instruction to transfer funds, combined with urgency, secrecy, and pressure to bypass normal payment controls.
- C. An email newsletter with stock photos.
- D. A chatbot that admits it's a bot.

Question 14 (True/False)

True or False: Responsible AI primarily means avoiding AI for any task that involves people.

- A. True
- B. False

Question 15 (MCQ)

A marketing team wants to use a striking AI-generated image that closely resembles a famous photographer's signature style and composition. The right next step before commercial use is:

- A. Publish quickly before anyone notices the resemblance.
- B. Route it through the normal originality, brand, and legal checks (e.g., reverse-image search, creative approval), and prefer alternatives if similarity is substantial.
- C. Add a small disclaimer that AI made it, which removes all copyright risk.
- D. Use it only on the company's largest billboard, where parody rules apply.

Question 16 (Situational Judgement)

Your team's AI tool assists with prioritizing incoming customer claims. Over a month you notice claims from certain postcodes consistently land at the bottom of the queue, with longer waits. Your most responsible action is:

- A. Ignore it — queue order is the tool's job.
- B. Manually fix the cases you personally notice and move on.
- C. Document the pattern and report it to the tool's owner/manager so the system can be investigated for proxy bias, while ensuring affected pending cases are reviewed.
- D. Post about the issue on social media to create pressure.

Question 17 (MCQ)

Why is “a human approves every AI decision” sometimes NOT genuine human oversight?

- A. Because humans are slower than machines.
- B. Because if the reviewer lacks the time, information, or authority to realistically disagree, the approval is a rubber stamp — oversight in form but not substance.
- C. Because human approval is illegal in automated systems.
- D. Because AI decisions cannot be reviewed once generated.

Question 18 (Multiple Select)

Which actions support transparency obligations? (Select all that apply.)

- A. Labelling a customer-service bot as a virtual assistant
- B. Disclosing significant AI assistance on a deliverable where company policy requires it
- C. Footnoting every use of spell-check in every email
- D. Following the organization's labelling rules for AI-generated images in published material
- E. Instructing a bot to deflect the question “am I talking to a human?”

Question 19 (MCQ)

A colleague says: “I didn't intend to mislead anyone — the AI added that fake case study, not me.” Which principle does this defense misunderstand?

- A. Tokens — the case study exceeded the token limit.
- B. Accountability — the person who reviews and sends content owns its claims; intent doesn't transfer responsibility to the tool.
- C. Fine-tuning — the model needed domain training.
- D. Context windows — the conversation was too long.

Question 20 (Situational Judgement)

You discover a teammate routinely uses AI to write personalized “hand-crafted” thank-you notes from your CEO to major clients, with the CEO unaware their name is being signed to them. The best characterization and action is:

- A. Harmless efficiency — clients enjoy the notes, so continue.
- B. Unauthorized impersonation — a real person's name and voice are being used without their knowledge; raise it so the practice is either authorized properly or stopped.
- C. A copyright problem only, to be solved by changing the wording.
- D. Acceptable as long as the AI tool is company-approved.

Quiz Answer Key & Explanations

Question 1 — Correct answer: B

Explanation: Models learn whatever patterns exist in their training data, including discriminatory ones. A numerical output is only as objective as the data and design behind it. The right move is investigating the skew against job-relevant criteria — not banning AI outright, and not assuming fairness.

Question 2 — Correct answer: C

Explanation: Voice cloning defeats recognition-based trust, and replying on the same channel just contacts the fraudster again. Urgency plus secrecy plus payment is the classic fraud signature: verify out-of-band, follow payment controls, and report — even if it turns out genuine.

Question 3 — Correct answer: A (True)

Explanation: Accountability does not transfer to software, and discrimination law generally concerns outcomes and practices, not just intentions. Deploying and relying on a tool is a human and organizational choice that carries responsibility for its effects.

Question 4 — Correct answer: B

Explanation: Identifying AI as AI to the people interacting with it is the core of transparency to recipients. A bot may be friendly and named, but actively denying being a bot is deliberate deception.

Question 5 — Correct answer: A, C, E

Explanation: Fabricated testimonials and unverified invented statistics are deception; unauthorized voice clones violate consent and likeness norms regardless of intent. Brainstorming questions and drafting reviewed emails are legitimate assisted-work uses.

Question 6 — Correct answer: B

Explanation: Consequential decisions about people must be explainable in human, job-relevant terms. An unexplained score means oversight existed on paper only. Hiding the score (D) worsens transparency rather than fixing the explainability gap.

Question 7 — Correct answer: B (False)

Explanation: AI output can resemble existing protected works closely enough to create risk, inputs you feed in may themselves be protected, and ownership of AI-generated work varies by jurisdiction and tool terms. Commercial use requires the same originality and rights checks as any other content.

Question 8 — Correct answer: C

Explanation: Records like minutes shape reputations; uneven representation is a fairness issue even when unintentional. The person circulating the document is accountable for how it represents people. Deleting the criticism (B) would distort the record further.

Question 9 — Correct answer: B

Explanation: Automation bias is a human factor: people defer to system outputs. It is why “a human approves every case” can be hollow if the human rubber-stamps — meaningful oversight needs time, information, and real authority to disagree.

Question 10 — Correct answer: B

Explanation: Candor is the disclosure test: uses of AI that depend on affected people not finding out are a red flag. It complements Permission (policy/data), Accuracy & fairness (verification/bias), and Takeover (human control of consequential decisions).

Question 11 — Correct answer: A, C, D

Explanation: Consent plus formal approval makes synthetic likenesses legitimate (A); reviewed drafting (C) and honestly labelled bots (D) are standard responsible patterns. Secret deepfakes (B) lack consent regardless of intent, and using a former employee's voice without authorization (E) violates likeness and consent norms.

Question 12 — Correct answer: C

Explanation: Authoritative-sounding statistics are a common hallucination, and once you send the memo, the claim is yours. Rewording (B) or pseudo-citing (D) doesn't create evidence. Decision-makers are entitled to know which claims rest on real sources.

Question 13 — Correct answer: B

Explanation: The fraud signature is behavioral, not technical: authority + urgency + secrecy + bypassing controls. Convincing audio/video no longer proves identity; out-of-band verification and fixed payment processes are the defense.

Question 14 — Correct answer: B (False)

Explanation: Responsible AI is deliberate use, not avoidance: matching tasks to AI strengths, keeping humans in charge of consequential decisions, checking fairness and accuracy, and being honest about AI's role. Blanket avoidance forgoes benefits without adding integrity.

Question 15 — Correct answer: B

Explanation: AI-generated content must clear the same rights and originality bar as any commercial asset. Disclaimers don't cure infringement, and prominence increases rather than reduces exposure.

Question 16 — Correct answer: C

Explanation: Geographic patterns can act as proxies for protected characteristics — a systemic fairness issue that individual workarounds (B) leave in place for everyone else. Escalating through proper channels gets the system examined; staying silent (A) makes you part of the failure.

Question 17 — Correct answer: B

Explanation: Meaningful oversight is a design property: reviewers need enough context, enough time per case, and real power to overrule. Automation bias plus impossible review volumes produces ceremonial oversight that protects no one.

Question 18 — Correct answer: A, B, D

Explanation: Transparency is proportionate and policy-guided: identify bots, disclose meaningful AI assistance where required, label generated media per policy. It does not demand absurd granularity (C), and deflecting identity questions (E) is concealment.

Question 19 — Correct answer: B

Explanation: Recipients experience the deception regardless of intent, and software cannot bear responsibility. Verification before sending is precisely the human step that accountability requires — skipping it is the failure, not the AI's addition.

Question 20 — Correct answer: B

Explanation: Sending messages in a real person's name without their knowledge is impersonation, even with positive intent and an approved tool. The fix is consent and process — the CEO authorizing a drafting workflow — not better wording. Tool approval governs the software, not the deception.

8. Final Assessment (Certification Style)

10 advanced scenario-based questions of medium difficulty, testing practical judgement across roles. Recommended passing score: 8 of 10. The answer key follows the assessment.

Question 1 (Scenario MCQ)

A hiring team adopts an AI shortlisting tool. Which combination of safeguards best reflects responsible deployment?

- A. Trust the vendor's "bias-free" claim and use the top-ranked candidates directly.
- B. Treat rankings as one input; have humans apply documented job-relevant criteria; monitor outcomes for skew across groups; require the vendor to evidence bias testing; keep decisions explainable to candidates.
- C. Use the tool but keep its existence secret from candidates and staff.
- D. Double the shortlist size so bias matters less.

Question 2 (Scenario MCQ)

Finance receives a video call appearing to include the CFO and two known directors, instructing an immediate confidential transfer. Everything looks authentic. The control that most reliably defeats this attack is:

- A. Training staff to spot visual glitches in deepfakes.
- B. Fixed payment processes that no call can override (dual authorization, callback verification on saved numbers, no exceptions for urgency or secrecy).
- C. Asking the callers a personal question during the call.
- D. Restricting video calls to approved platforms.

Question 3 (Scenario MCQ)

A support chatbot resolves issues well, but transcripts show it answers "Are you a real person?" with "I'm here to help you today!". The appropriate fix is:

- A. Nothing — the deflection is technically not a lie.
- B. Configure the bot to state clearly that it is a virtual assistant and to offer a path to a human agent.
- C. Remove the customers' ability to ask the question.
- D. Make the bot's avatar more obviously robotic but keep the deflection.

Question 4 (Scenario MCQ)

Marketing wants to launch a campaign with an AI-generated jingle that testers say "sounds just like" a current chart hit, and AI images echoing a famous artist's style. The soundest position is:

- A. Proceed — AI-generated means legally original.
- B. Proceed but credit “AI” as the artist.
- C. Pause both assets for rights/originality review; substantial similarity to existing works creates infringement and reputational risk that disclaimers don't cure; commission more original directions if needed.
- D. Use the assets only in markets where the original artists are less famous.

Question 5 (Scenario MCQ)

An operations AI flags transactions for manual fraud review. Reviewers handle 400 flags/day each and approve the AI's recommendation 99.8% of the time, averaging 40 seconds per case. Leadership cites this as “strong human oversight.” The most accurate critique is:

- A. Oversight is strong because every case sees a human.
- B. The review volume and near-total agreement suggest rubber-stamping driven by automation bias; meaningful oversight needs realistic caseloads, adequate information, and demonstrated willingness/authority to overrule.
- C. The AI should review the humans instead.
- D. Oversight is unnecessary once accuracy exceeds 99%.

Question 6 (Scenario MCQ)

A junior analyst privately tells you they were instructed by a team lead to generate AI “customer success stories” for a bid, based on no real customers. Your most responsible course is:

- A. Advise the analyst to comply — refusing could hurt their career.
- B. Help make the fake stories more plausible so the deception is at least lower-risk.
- C. Support the analyst in declining, and escalate the instruction through the appropriate channel (manager, compliance, or ethics reporting), since fabricated evidence in a bid is deception with legal exposure.
- D. Suggest the analyst quietly swap in stories from a previous employer.

Question 7 (Scenario MCQ)

An employee challenges their low AI-assisted performance summary, noting it omits two major achievements documented in the project tracker. The manager realizes the AI summarized only email threads it was given. The principled resolution is:

- A. Keep the summary — re-running the process would undermine the tool's credibility.
- B. Correct the record using the full evidence, ensure the employee's documented achievements are reflected, and fix the process so source coverage is checked before summaries influence evaluations.
- C. Tell the employee to email more so future summaries capture them.

D. Average the AI summary with the employee's self-assessment.

Question 8 (Scenario MCQ)

Your organization is drafting guidance on synthetic media of real people. Which rule best captures the responsible-AI position?

- A. Allowed for internal humor, banned externally.
- B. Allowed only with the depicted person's informed consent and formal policy approval, with required labelling — no exceptions for positive intent or internal audiences.
- C. Banned entirely, including consented, approved uses such as training avatars.
- D. Allowed if the media is clearly higher quality than typical deepfakes.

Question 9 (Scenario MCQ)

A planner notices that AI-drafted vendor communications are consistently warmer and more flexible toward vendors with Western-sounding names. The pattern is subtle but repeats. This is best understood as:

- A. Coincidence; tone has no fairness dimension.
- B. A potential learned bias surfacing in everyday output — worth documenting, correcting in drafts before sending, and reporting to the tool owner for investigation.
- C. A reason to stop using names in any business communication.
- D. Acceptable, since vendors are companies rather than protected individuals.

Question 10 (Scenario MCQ)

Leadership asks for one operating rule that best integrates this module. The strongest candidate is:

- A. "AI may draft anything, and whoever sends it owns it: verify claims, check fairness to people, disclose honestly, and keep humans genuinely deciding the consequential outcomes."
- B. "AI use is acceptable whenever it saves more than an hour."
- C. "Ethics review applies only to customer-facing AI."
- D. "Any output from an approved tool is by definition responsible."

Final Assessment Answer Key

Question 1 — Correct answer: B

Explanation: Option B combines human decision rights, explainability, outcome monitoring, and vendor accountability — the full responsible pattern. Vendor claims aren't evidence (A), secrecy violates transparency (C), and enlarging a skewed list doesn't correct the skew (D).

Question 2 — Correct answer: B

Explanation: Detection skills decay as fakes improve, and fraudsters research personal details. Procedural controls are robust precisely because they don't depend on perceiving the fake: legitimate requests survive a callback and dual authorization; fraudulent ones don't.

Question 3 — Correct answer: B

Explanation: Evasive non-answers to direct identity questions are concealment, which transparency norms (and some regulations) prohibit. Honest identification plus easy escalation preserves both trust and the bot's usefulness.

Question 4 — Correct answer: C

Explanation: Similarity, not the production method, drives copyright risk. Crediting AI or choosing different markets does not address infringement. Standard creative-rights review applies to generated content exactly as to agency work.

Question 5 — Correct answer: B

Explanation: Oversight is measured by its substance, not its presence on a flowchart. Forty seconds per case at that volume makes genuine scrutiny implausible; the agreement rate is a symptom, not proof of model quality. Stakes-appropriate review design is the fix.

Question 6 — Correct answer: C

Explanation: Fabricated success stories in a competitive bid are fraud-adjacent deception; “following instructions” doesn't dissolve individual responsibility. Improving (B) or recycling confidential material (D) compounds the violation. Escalation channels exist precisely for instructed misconduct.

Question 7 — Correct answer: B

Explanation: Fairness and accuracy in records about people outrank tool credibility. The root cause is incomplete context (source coverage), a process flaw to fix — not the employee's communication style. Averaging (D) institutionalizes a known error.

Question 8 — Correct answer: B

Explanation: Consent plus governance is the dividing line that permits legitimate uses (approved avatars) while excluding secret “jokes” and impersonation. Internal/external and quality distinctions miss the point; a total ban (C) needlessly forbids consented, governed uses.

Question 9 — Correct answer: B

Explanation: Bias manifests in tone and treatment, not only in scores — and name-correlated differences are a recognized proxy pattern. The dual response is individual (fix drafts you send) and systemic (report so the tool is examined). Fairness norms extend to business conduct generally.

Question 10 — Correct answer: A

Explanation: Option A unifies accountability, verification, fairness, transparency, and meaningful oversight into one usable rule. Time saved (B) is irrelevant to ethics; internal uses also affect people (C); and tool approval governs software choice, never the ethics of a particular use (D).